

CS4225/CS5425 Big Data Systems for Data Science

Revision (1st Half)

Bingsheng He
School of Computing
National University of Singapore
hebs@comp.nus.edu.sg



Exam Info

- Time: 29 APR 2026 09:00-11:00
 - Students are allowed to enter the venue 10 mins before exam starts.
 - Students will not be permitted to enter the venue after 1 hour from the start of the exam.
- Venue: TBD
- F2F - Hardcopy (pen and paper)
- Open Book Exam
 - Any physical materials are allowed
 - Calculator is allowed
 - Any other electronics devices are NOT allowed

Exam



- Focus is on understanding and application, not facts / memorization
- Question structures (total 50 marks):
 - True / False question with a short explanation / justification
 - Application / Scenario Based Question
 - Give you a practical scenario and let you come out a solution / suggestion
- Example questions
 - **Integrative**: Require you to combine knowledge from different weeks of content
 - **“Application”**: Require you to apply your knowledge of fundamental concepts to reasonably practical scenarios.
 - **“Why not”**: Example, Tommy proposed a solution A to solve problem B. Tell me what is the problem with solution A and how to overcome this problem

Scope of Exam

- **Scope:** content discussed in the lecture which appears in slides
- **Out of scope:**
 - Content only found in the notes underneath the slides (these can be ignored)
 - Content only discussed in response to student questions (i.e. beyond the scope of the slides)
 - Content marked with “Details” or that I mentioned is out of scope
 - If there are any tasks which ask for code, they will allow pseudocode. As long as your pseudocode is reasonable / understandable by the grader, we will accept it.
 - Historical details
 - Further readings
 - Supplementary slides



Infrastructure for Big Data

- Data center architecture
- Concepts: capacity, latency, bandwidth, throughput
 - General / conceptual understanding is sufficient, no need exact values
- Storage hierarchy in a data center
- Capacity / latency / bandwidth over a data center (bottom right plot)
- Big Ideas (scale out, move processing to data, sequential processing, scalability)




MapReduce

- Design principles: motivation, challenges, key ideas
- MapReduce
 - MapReduce interface
 - Ensuring correct types when writing MapReduce programs
 - Execution framework: responsibilities of Hadoop
- MapReduce framework implementation
 - Understand map, shuffle, reduce phases (no need for implementation details beyond the level discussed in lecture)
 - Understand how work is divided; difference between workers, tasks, functions, and tuples





MapReduce

- Role of partitioner and combiner
 - Benefit of partitioners / combiners; correct usage of combiners
- Writing MapReduce programs
- Performance analysis: disk and network I/O, and working set memory usage



MapReduce


- Preserving State in MapReduce tasks: how to use it, what it can be used for
- [Optional] Secondary Sort / Composite Keys: how to use it, what it can be used for
- Regular combiners vs in-mapper combiners: usage, efficiency gains



```
class Mapper {  
    val counts = new Map()  
  
    def map(key: Long, value: Text) = {  
        for (word <- tokenize(value)) {  
            counts(word) += 1  
        }  
    }  
  
    def cleanup() = {  
        for ((k, v) <- counts) {  
            emit(k, v)  
        }  
    }  
}
```

HDFS




- Basic motivation and interface (no need to be familiar with specific commands)
- Basic infrastructure (namenode / datanode, reading / writing, replication)



MapReduce and Relational Databases

- Understanding of algorithms and MapReduce implementation


- Basic operations:
Projection, Selection,
Group By
- Joins:
 - Broadcast Join
 - Reduce Side Join



MapReduce and Data Mining


○ Similarity Search

- Similarity measures
- Goals: similarity search vs all-pairs similarity
- Shingling
- MinHash algorithm & key property (no need for proof)
- MapReduce implementation




○ K-means clustering

- K-means algorithm
- MapReduce implementation
- Efficiency analysis




NoSQL

- Types of NoSQL databases
- Types of data suitable for key-value, document, graph, vector databases
- [Optional] Concepts: eventual consistency, duplication



All readers read new value (60kg)



Eventually, the update is propagated to Node C; after which, readers will read the correct value

Strong consistency: any reads immediately after an update must give the same result on all observers

Eventual consistency: if the system is functioning and we wait long enough, eventually all reads will return the last written value

Tips for Designing MapReduce Programs

- In exam, only pseudocode will be required (if you want to write in Python / Java, that is fine too). As long as your pseudocode is “reasonable” (understandable to the grader), it will be accepted
- Think about what key the mapper should emit: this key will be used to group the data for the reducers
- Then think about what value the mapper should emit: this should give the reducers all the information they need to perform their task
- Evaluating efficiency: main considerations are the disk and network I/O (determined by the amount of data emitted by mappers, but possibly reduced by combiners), and the memory working set (determined by the mapper’s intermediate state)

MapReduce: Example

Given two documents (documents 0 and 1), we want to find all k-shingles in document 0 but not in document 1 (that is, all shingles which appear at least once in document 0, but do not appear at all in document 1), where $k=1$. We receive our documents line by line, receiving input key-value pairs of the form `<docID, line>`, where `docID` (an integer of either 0 or 1) is the document ID being read, and `line` is a space-separated string containing a line of the document. For each k-shingle in the final answer, your reduce function should emit a tuple (only once) of the form `<shingle, 1>`. Show pseudo-code for how you would use MapReduce to find all such shingles. You can assume that the input (in `line`) is ‘clean’; e.g. no duplicate spaces or spaces at the start / end of the line, and no characters other than letters and spaces are present. You can assume the existence of a string splitting function of your choice, e.g. `split()`.

Bonus: show how to use a combiner (not in-mapper combiner) to speed up the program.

```
map (docID, line) {  
    /* your pseudo code */  
    /* output the map results by calling the API, emit(specify your map output) */  
}  
  
reduce /* specify your input to reducer */ {  
    /* your pseudo code */  
    /* output the map results by calling the API, emit(shingle, 1). */  
}
```

MapReduce: Example

```
map (docID, line) {  
    tokens = line.split()  
    for token in tokens:  
        emit(token, docID)  
}  
  
reduce (token, docIDs[]) {  
    if 0 in docIDs and 1 not in docIDs:  
        emit(token, 1)  
}
```

Explanation:

- 1-shingles are just words (or tokens)
- The map() function should use words as keys, so that each reduce() will handle one word
- The map() needs to emit the docID as value, since the reduce() needs this information

MapReduce: Example

```
map (docID, line) {  
    tokens = line.split()  
    for token in tokens:  
        emit(token, docID)  
}  
  
reduce (token, docIDs[]) {  
    if 0 in docIDs and 1 not in docIDs:  
        emit(token, 1)  
}  
  
combine (token, docIDs[]) {  
    if 0 in docIDs:  
        emit(token, 0)  
    if 1 in docIDs:  
        emit(token, 1)  
}
```

Explanation:

- 1-shingles are just words (or tokens)
- The map() function should use words as keys, so that each reduce() will handle one word
- The map() needs to emit the docID as value, since the reduce() needs this information

Combiner:

- The combine() function combines multiple (token, 0) tuples into a single (token, 0) tuple, and similarly combines multiple (token, 1) tuples into a single (token, 1) tuple.
- Thus it does not change the final output, but speeds up the program by reducing the number of such tuples emitted (and thus the disk and network I/O).



**Good luck
with the
exam!**

Acknowledgement and Others

- Some slides are adopted/revised from
 - Bryan Hooi
- Time management is the key.
 - Effective Time Management: Best Practices for Students and Educators. <https://blog.nus.edu.sg/learners/2024/09/24/effective-time-management-best-practices-for-students-and-educators/>.
 - 8 Time Management Tips for Students.
<https://summer.harvard.edu/blog/8-time-management-tips-for-students/>.